

# Fine tuning a VGG-16

## Motivation

- There are "famous" network architectures like AlexNet, GoogLeNet, NetworkInNetwork, VGG, Inception, ResNet, FCN8, SegNet etc.
- because they are top performing in some task on some benchmark dataset like classification of ImageNet, ILSVRC-2012, semantic segmentation of Cityscapes etc.
- Authors have made them publicly available in the framework they used to develop (Keras, Tensorflow, Theano, Lassagne, Caffe... many!)
- What's available is not only the architecture and the source code creating it, but also the weights (layers parameters) after a long and careful training with a possibly huge dataset
- There's even something called a zoo : a repository of models and their weights, for instance in Caffe (go [here \(https://github.com/BVLC/caffe/wiki/Model-Zoo\)](https://github.com/BVLC/caffe/wiki/Model-Zoo))
- Maybe someone has managed to translate one such model from one framework to yours (Tensorflow), or if not, there are tools to do it like [Caffe-to-Tensor \(https://github.com/ethereon/caffe-tensorflow\)](https://github.com/ethereon/caffe-tensorflow)

## What's fine tuning a network

You'll like to reuse one such model for

- the **same type of task**, say, image classification, **but**
- on a **different type of images** : instead of ImageNet = 1.4M images of 1000 classes of objects, a dataset of images of cars where classes are make + model

Instead of start training the model from scratch, since the task is the same you want to perform **domain transfer** : adapt some of the net parameters to the new images.


Why ?

- the original net is large  $\Rightarrow$  hard to train
- already performs well on another difficult dataset
- filters of lower layers are probably ok
- so we only want to adapt the upper conv layers or just the classifier layers

**Fine tuning** is taking the parameters of one network as the starting point and retrain (usually just some layers) with new images of a different domain.

## VGG-16

VGG achieves 92.7% top-5 test accuracy in ImageNet. We'll retrain VGG-16 to classify cars.



Here (<https://github.com/machrisaa/tensorflow-vgg>) there's a Python class Vgg16 that instantiates the VGG-16 model and loads all the weights.

We'll change it to

- accept  $112 \times 112 \times 3$  images instead of  $224 \times 224 \times 3$
- just load the first 5 convolutional layers
- and not the two fully-connected ones that perform the classification ( $7 \cdot 7 \cdot 512 = 25088$  to  $4096 + 4096$  to  $1000$  classes =  $120\text{M}$  params.)
- replace the fully connected layers by 2 smaller fully connected layers ( $4 \cdot 4 \cdot 512$  to  $512 + 512$  to  $41$  classes =  $4\text{M}$  params.)
- train just these 2 new layers

In [ ]:

```
import os
import sys

import numpy as np
import tensorflow as tf

VGG_MEAN = [103.939, 116.779, 123.68]

class My_Vgg16:
    def __init__(self, vgg16_npy_path=None):
        if vgg16_npy_path is None:
            path = sys.modules[self.__class__.__module__].__file__
            # print path
            path = os.path.abspath(os.path.join(path, os.pardir))
            # print path
            path = os.path.join(path, "vgg16.npy")
            print(path)
            vgg16_npy_path = path

        self.data_dict = np.load(vgg16_npy_path, encoding='latin1').item()
        print("np file loaded")

    def _max_pool(self, bottom, name):
        return tf.nn.max_pool(bottom, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1],
                               padding='SAME', name=name)

    def _conv_layer(self, bottom, name):
        with tf.variable_scope(name) as scope:
            filt = self.get_conv_filter(name)
            conv = tf.nn.conv2d(bottom, filt, [1, 1, 1, 1], padding='SAME')

            conv_biases = self.get_bias(name)
            bias = tf.nn.bias_add(conv, conv_biases)

            relu = tf.nn.relu(bias)
            return relu

    def get_conv_filter(self, name):
        return tf.Variable(self.data_dict[name][0], name="filter")
        # use tf.constant() to prevent retraining it accidentally

    def get_bias(self, name):
        return tf.Variable(self.data_dict[name][1], name="biases")
        # use tf.constant() to prevent retraining it accidentally

    def build(self, rgb, train=False):
        rgb_scaled = rgb * 255.0

        # Convert RGB to BGR
        red, green, blue = tf.split(3, 3, rgb_scaled)
        assert red.get_shape().as_list()[1:] == [112, 112, 1]
        assert green.get_shape().as_list()[1:] == [112, 112, 1]
        assert blue.get_shape().as_list()[1:] == [112, 112, 1]
```

```

bgr = tf.concat(3, [
    blue - VGG_MEAN[0],
    green - VGG_MEAN[1],
    red - VGG_MEAN[2],
])
assert bgr.get_shape().as_list()[1:] == [112, 112, 3]

self.conv1_1 = self._conv_layer(bgr, "conv1_1")
self.conv1_2 = self._conv_layer(self.conv1_1, "conv1_2")
self.pool1 = self._max_pool(self.conv1_2, 'pool1')

self.conv2_1 = self._conv_layer(self.pool1, "conv2_1")
self.conv2_2 = self._conv_layer(self.conv2_1, "conv2_2")
self.pool2 = self._max_pool(self.conv2_2, 'pool2')

self.conv3_1 = self._conv_layer(self.pool2, "conv3_1")
self.conv3_2 = self._conv_layer(self.conv3_1, "conv3_2")
self.conv3_3 = self._conv_layer(self.conv3_2, "conv3_3")
self.pool3 = self._max_pool(self.conv3_3, 'pool3')

self.conv4_1 = self._conv_layer(self.pool3, "conv4_1")
self.conv4_2 = self._conv_layer(self.conv4_1, "conv4_2")
self.conv4_3 = self._conv_layer(self.conv4_2, "conv4_3")
self.pool4 = self._max_pool(self.conv4_3, 'pool4')

self.conv5_1 = self._conv_layer(self.pool4, "conv5_1")
self.conv5_2 = self._conv_layer(self.conv5_1, "conv5_2")
self.conv5_3 = self._conv_layer(self.conv5_2, "conv5_3")
self.pool5 = self._max_pool(self.conv5_3, 'pool5')

shape_pool5 = self.pool5.get_shape().as_list()[1:]
assert 'pool5', shape_pool5 == [4, 4, 512]

"""
The new fully connected layers
"""
dim1 = np.prod(shape_pool5) # 4*4*512 = 8192
dim2 = 512
n_output = 41 # number of classes = car makes
fc_weights = {
    'wd1': tf.Variable(tf.random_normal([dim1, dim2], stddev=0.1), name=
'wd1'),
    'wd2': tf.Variable(tf.random_normal([dim2, n_output], stddev=0.1), n
ame='wd2')
}
fc_biases = {
    'bd1': tf.Variable(tf.random_normal([dim2], stddev=0.1), name='bd1'
),
    'bd2': tf.Variable(tf.random_normal([n_output], stddev=0.1), name='b
d2')
}

x = tf.reshape(self.pool5, [-1, dim1], name="flat_pool5")
self.fc6 = tf.nn.bias_add(tf.matmul(x, fc_weights['wd1']), fc_biases['bd
1'], name="fc6")
self.relu6 = tf.nn.relu(self.fc6)
if train:
    self.relu6 = tf.nn.dropout(self.relu6, 0.75)

# the two outputs, logits is for training, probs for testing
self.logits = tf.nn.bias_add(tf.matmul(self.relu6, fc_weights['wd2']), f

```

```
c_biases['bd2'], name="logits")
    self.probs = tf.nn.softmax(self.logits, name="probs")
```

## The dataset

It's the same one we used in the Feeding and Queues section.

In [ ]:

```
from feeding_and_queues.dataset import Dataset
```

## Run fine tuning

It seems that it's better to reduce the original learning rate, something like

- 1/10th for top layers
- 1/100 for intermediate layers

In [ ]:

```
import numpy as np
import time
import platform
import matplotlib.pyplot as plt

from skimage.io import imread

import tensorflow as tf
from fine_tuning_vgg.my_vgg16 import My_Vgg16
from feeding_and_queues.dataset import Dataset

tf.reset_default_graph()

xs = tf.placeholder(tf.float32, [None, 112, 112, 3])
ys = tf.placeholder(tf.float32, [None])
# the VGG-16 weights are here: https://dl.dropboxusercontent.com/u/50333326/vgg16.npy
# Based on this implementation https://github.com/ry/tensorflow-vgg16
if platform.node().upper()=='CVC220':
    vgg16_npy_path = '/home/joans/Documents/recearca/tensorflow/implementacions/tensorflow_vgg_master/vgg16.npy'
elif platform.node().upper()=='CVC180':
    vgg16_npy_path = '/home/joans/Documents/classe_tf/notebooks/fine_tuning_vgg/vgg16.npy'
else:
    raise Exception('Not ready to run on computer '+platform.node())

vgg = My_Vgg16(vgg16_npy_path=vgg16_npy_path)
with tf.name_scope("my_vgg16"):
    vgg.build(xs)

print '\ntrainable variables'
for v in tf.trainable_variables():
    print v.name, v.get_shape().as_list()

"""
We want to train only the weights and biases of the two
fully connected layers.
"""

vars_to_optimize = [v for v in tf.trainable_variables() \
    if v.name.startswith('my_vgg16/wd') \
    or v.name.startswith('my_vgg16/bd')]

print '\nvariables to optimize'
for v in vars_to_optimize:
    print v.name, v.get_shape().as_list()

learning_rate = 0.001/10.
loss = tf.reduce_mean(
    tf.nn.sparse_softmax_cross_entropy_with_logits(vgg.logits, tf.to_int64(ys)))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)
"""
minimize with list of variables to update
"""
train_op = optimizer.minimize(loss, var_list=vars_to_optimize)

corr = tf.equal(tf.argmax(vgg.probs, 1), tf.to_int64(ys)) # count corrects
```

```

accr = tf.reduce_mean(tf.cast(corr, tf.float32)) # accuracy
init = tf.global_variables_initializer()

ds = Dataset()
ds.new_height, ds.new_width = (112, 112)
batch_size = 10
num_epochs = 5

with tf.Session() as sess:
    sess.run(init)
    ds.epochs_completed
    while ds.epochs_completed < num_epochs:
        ds.epochs_completed
        batch_xs, batch_ys = ds.next_batch(batch_size)
        # compute average loss and accuracy for each batch
        _, batch_cost, batch_acc = sess.run([train_op, loss, accr],
                                             feed_dict={xs: batch_xs, ys: batch_ys})
    print("cost: %.9f, acc %.9f" % (batch_cost, batch_acc))

```